

CZECH TECHNICAL UNIVERSITY
FACULTY OF NUCLEAR SCIENCES AND PHYSICAL
ENGINEERING

Department of Software Engineering
Field: Engineering Informatics



Navigation in virtual environment Navigování ve virtuálním světě

MASTER THESIS

Author: Bc. Václav Honzík
Supervisor: Doc. Ing. Zdeněk Žabokrtský, Ph.D.
Year: 2014

Insert assignment instead of this page before handing in the thesis!!!!

Declaration

I declare, that I have developed my master thesis independently and used only the materials (literature, projects, SW etc.) listed in attached list.

In Prague on

.....
Bc. Václav Honzík

Acknowledgment

I would like to thank Doc. Ing. Zdeněk Žabokrtský, Ph.D. for supervising and shaping my master thesis. I would like to also thank Prof. Kristina Striegnitz, Ph.D. for giving me the opportunity and helping me with my master thesis.

Bc. Václav Honzík

Title:

Navigation in virtual environment

Author: Bc. Václav Honzík

Field: Engineering Informatics

Type of thesis: Master thesis

Supervisor: Doc. Ing. Zdeněk Žabokrtský, Ph.D.
Institute of Formal and Applied Linguistic, Faculty of Mathematics
and Physics, Charles University in Prague

Advisor: Prof. Kristina Striegnitz, Ph.D.
Union College, Schenectady, NY USA

Abstract: This thesis discusses a natural language generation in a virtual environment navigational task. It focuses on interesting referencing strategies which was discovered in available dataset of human to human interaction. Because these strategies does not follow common methodology of referring expression generation field, this paper analyses them and attempts to model them using machine learning techniques. It also presents an experiment assessing their effect on the task proficiency. The results suggest that such strategies are too complex to model using spatial information from the scene and human factors must be taken into account. The experiment concludes that they FILL RESULTS effect on the task proficiency.

Key words: navigation, referring expressions, language realization

Název práce:

Navigování ve virtuálním světě

Autor: Bc. Václav Honzík

Abstrakt: Tato práce se zabývá generováním přirozeného jazyka pro navigování ve virtuálním světě. Předmětem zájmu jsou zajímavé strategie odkazování, které byly objeveny v dostupných datech v nichž jedna osoba navigovala druhou. Tyto strategie se vymykají běžným metodám v disciplíně generování odkazujících výrazů a tak je tato práce analyzuje a pokouší se je modelovat technikami strojového učení. V rámci této práce je také proveden experiment studující vliv používání těchto strategií na efektivnost navigování. Výsledky naznačují, že tyto strategie jsou příliš komplexní na to aby mohly být úspěšně modelovány prostorovými informacemi ze scény a musí k tomu být využity i lidské faktory. Experiment ukázal, že strategie DOLPNIT vliv na efektivnost navigování.

Klíčová slova: navigace, odkazující výrazy, jazyková realizace

Contents

Introduction	8
1 Background	9
1.1 Referring expression generation: introduction and historical overview	10
1.2 Referring expression generation: current research	11
2 GIVE Challenge	14
2.1 Introduction	14
2.2 History of GIVE Challenge	16
2.3 Task and GIVE world	18
3 The S-GIVE Dataset	21
3.1 General overview	21
3.2 World and Demographic factors	27
3.3 Referring expressions	30
3.4 Chains of references	32
4 Machine Learning on REs	35
4.1 Timing of the first reference	35
4.2 Chains of references	37
4.2.1 Presence of the chains	39
4.2.2 Chain length prediction	41
4.2.3 Closer look at chains' content	42
4.3 Room memory	45
4.4 Thoughts on ML in S-GIVE dataset	46
5 Experiment on RE	47

5.1	Hypothesis	47
5.2	Experimental set-up	48
5.3	Results	52
Conclusion		54
Appendix		58
.1	Histograms for first reference ML	59
.2	Histograms for chains ML	65

Introduction

This thesis deals with language generation for navigating in a virtual environment. A dataset of spoken interactions from a shared navigation task is used to model the language generation. The thesis explores a particular strategy of referring, which differs from the traditional methodology used in the relevant research.

It consists of five chapters.

The first chapter describes the background of the navigation and especially of the language generation in navigation. It introduces the domain of Referring Expression Generation and a related research.

In the second chapter the GIVE framework is introduced. The whole thesis is built upon this framework and is therefore vital for the following chapters.

The third chapter analyses the spoken interaction of the S-GIVE dataset. The S-GIVE dataset was created in the GIVE framework and this thesis is one of the first papers to analyse it.

The fourth chapter attempts to create models of Natural Language Generation using machine learning techniques. Using the data from the S-GIVE dataset, it explores various sub-problems of language generation.

In the last chapter an experiment to evaluate the mentioned referring strategy and its results are presented. Once again built upon the GIVE framework, it compares the newly discovered strategy with more traditional approaches.

Chapter 1

Background

Navigation is a broad problem involving different scientific areas. Generally speaking, a path to the destination must be found first. That problem is solved by path-planning, a problem commonly associated with robotics, and specifically by pathfinding, a well-known and studied problem of artificial intelligence. Once a path is known, it must be realized into a natural language. The realization is studied by Natural Language Generation (NLG), a part of the discipline of Natural Language Processing (NLP), which itself belongs to Computational Linguistics (CL). On top of these two subproblems, one has to consider the domain the navigation is intended for. Differences in the design between car navigation systems and art gallery guides are immense.

I would like to briefly touch on why navigation is still an interesting problem, when navigation devices are nowadays almost a necessary part of the car equipment. Car navigation is a specific navigation problem. Great availability of maps is one of the reasons why car navigation is so developed and massively produced. It has limited sense of direction, because cars only move on roads, as far as navigation is concerned anyway. The roads can fork and form complicated cross-roads, but that isn't still nowhere close to the free movement of, for example, walking. The limited sense of directions ties to a relatively limited vocabulary. Thanks to a usually well-defined space where car moves, the navigation system's world representation can be relatively straightforward. Car navigation is still a complicated problem, but being able to relatively well navigate drivers around mapped areas tells us little about other navigation in other domains.

This thesis studies navigation in virtual environment when the navigated person is walking in a house-like environment and apart from moving around must perform other tasks. This environment is radically different from the problem of car navigation.

The primary interest of this thesis is NLG part of navigation and more specifically referring expression generation, which I will introduce in the first section.

The second section summarizes related research in referring expression generation.

1.1 Referring expression generation: introduction and historical overview

From the complex and comprehensive problem of navigating persons, I'm especially interested in the language generation subproblem. Moreover, I limited my research mostly to the domain of Referring Expression Generation (REG). REG is a part of NLG. In the dataset and framework, which I will present in the following chapters, REG is an especially important part of language generation. In fact, in the virtual world which I'm working with, extending REG system to a complete language generation system is a trivial task of adding few verbs.

Krahmer and Van Deemter (2012) created a very well written survey of history and development of REG up to recent times. Following the principle of "Don't reinvent the wheel," I will provide only a short explanation of what REG is and a brief overview of its history both of them being heavily inspired by the survey.

REG belongs to the domain of Natural Language Generation (NLG). It is concerned with generating referring expressions (RE) to an object or objects of interest. Suppose we have three buttons next to each other and we need one of them to be pressed. Also suppose we are not able to do it ourself right now, but another person is available nearby. Most individuals would have no trouble to address the person nearby and ask him to press the button. Part of their utterance would "point out" which button of the three needs to be pressed. That part of the utterance is a RE. Producing understandable and effective REs is for most speakers relatively easy task. But for computer programs it is not so. The context of the real world application is usually very large and if we take into account mutual relationships between entities in the context (such as a button is next to a painting), the number of possible combinations quickly grows to problematic numbers.

The first REG research appeared in the 1980s. Krahmer and Van Deemter (2012) state that, influenced by methodologies of computation linguistic at that time, they studied REG as a part of larger speech acts and doing so on hard and often anomalous cases. In 1990s, a famous paper by Dale and Reiter (1995) shifted the focus to determining which properties should be used, when the goal is to identify the referent, while avoiding being more informative than is required. The new aim of REG was generating human-like descriptions. According to Krahmer and Van Deemter (2012), the 1990s also spawned first REG algorithms for well defined REG problem, such as the influential Incremental Algorithm. However, the research was limited to the target being just one object, simple knowledge representation, no vague properties, all objects being equally salient and ignoring the stage of surface realization of the chosen properties. "A substantial part of recent REG research is dedicated to lifting one or more of these simplifying assumptions." (Krahmer and Van Deemter, 2012)

Apart from lifting these assumptions, recent REG research has been interested in the evaluation of REG algorithms. We can also see a tendency to move from simple well-defined environments, towards more natural and complex ones, such as the one this thesis is examining.

Exploring a dataset of spoken navigation through a complex 3D virtual world, I have noticed some behaviours which slightly deviates from the conventional focus of REG research. Speakers do not necessarily produce a reference which uniquely identifies the referent in the current context. Instead, they produce a reference which only partially identifies the referent first and then rely on feedback and additional REs when necessary to identify the referent. This thesis is primarily focused on this strategy.

Having introduced the REG, I'll now move onto related research.

1.2 Referring expression generation: current research

In this section I will briefly present work from the area of REG, which I deem relevant to my research.

Ha et al. (2012) talk about an ‘information gap’ caused by existence of a non-dialogue communication stream. They concluded that the posture of user, an example of implicit information from the non-dialogue streams, is a significant attribute in the modeling of dialog acts. Their goal is to overcome this ‘information gap’ through machine learning techniques. A shared view of the virtual world in this thesis is also a form of non-dialogue stream, with which the navigation system must work. I also try to apply machine learning to help with language generation.

Viethen et al. (2011a) compare traditional algorithmic approaches with alignment approaches based on psycho-linguistic models of REG. They use a large data-set (16,358 referring expressions) of a direction giving task on a shared 2D visual scene introduced by Louwerse et al. (2007). They use three feature sets: traditional REG set, alignment set and independent set (general information about the scene) to build decision tree models (concretely C4.5) to predict content patterns in subsequent references. The traditional REG set includes features such as the distance to the closest visual distractor and the number of visual distractors. The alignment set includes features such as distance in REs to the last use of the predicted attribute for the target and how often has the attribute been used for the target. Most frequent type of landmark is one of the features in the independent set. The alignment based models outperformed the traditional REG ones and the best model combines all feature sets to achieve an accuracy of 58.8% and a DICE score of 0.81. Not using traditional algorithmic REG features did not result in a significant decrease of accuracy, suggesting that the visual context doesn't play such an important role as it was believed in the REG research so far. Viethen et al. (2011b) verified this surprising conclusion by varying the visual context. They argue that the relative simplicity of visual scenes used in contemporary research might be the cause of insignificance of the visual context. I would argue that the 3D virtual world explored in this thesis is more complex then theirs and therefore this paper can provide some further insight into these questions.

Stoia et al. (2006a) were interested in the timing of the first reference to a target in a 3D virtual world. They predicted whether a direction giver refers to the target

or delays the reference based on the spatial data. Their attributes were angle and distance to the target, the number of visible distractors (either only these with the same category as the target or all of them) and whether the target is visible. The most important feature in the decision tree model was the number of visible distractors followed by angle and distance. They achieved 86% accuracy, compared to a 70% baseline. The baseline was to refer when the target is visible and to delay the reference when the target isn't visible. Part of the machine learning attempts of this thesis replicate the first reference timing experiment of Stoia et al. (2006a) on the S-GIVE dataset.

Stoia et al. (2006b) developed decision trees to generate a noun phrase, specified by three slots: determiner/quantifier, pre-modifier/post-modifier and head noun. They used a data-set from a 3D virtual world navigation task similar to the S-GIVE dataset. Four categories of features were used: dialog history, spatial and visual features, relation to other objects in the world and the object category. The decision trees revealed significant dependencies between the slots and the importance of the spatial features. Interestingly, they used three types of system evaluation. The exact match evaluation produced 31.2% accuracy compared to a 20% most-frequent baseline. A comparison with a hand-crafted Centering algorithm (Kibble and Power, 2000) ended with similar accuracy, favoring the machine-learning approach for requiring less structural analysis of the input text. Lastly, when humans judged the system output, it was at least equal or preferred to the original spontaneous language in 62.6% (inter-annotator reliability $\kappa = 0.51$).

Gallo et al. (2008) showed that the Fruit Carts corpus can be used in NLG by case study on message complexity and structural realizations. In Fruit Carts corpus one participant - the speaker - is given a map with geometric shapes and fruits in a specific configuration. The speaker's task is to instruct the second participant - the actor - to reorganize the objects in the actor's scene to match the speaker's map. A logistic regression confirmed that the complexity of verb arguments affects production choice between mono-clausal or bi-clausal structure. In other words, people tend to separate complex instructions into multiple clauses. Therefore the complexity of the virtual environment affects how people speak on all linguistic levels. REG should take that into consideration.

Clark and Krych (2004) examined speakers' monitoring of addressees in a Lego-building experiment. One participant - the director - knew 10 Lego models and how to build them. The director was verbally instructing the second participant - the builder - to build these models. In one group the director could see the builders workspace, in a second group he could not and in a third the instructions were audio-taped and simply passed onto the builder. Builders communicated with the directors on the workspace through head gestures and manipulating blocks (placing, exhibiting or posing and so on). When the workspace was blocked off, the task took much longer. In the audio-taped group the builders made many more errors. Directors often altered their utterances mid-course based on the builders actions. This ties to the importance of a shared visual context and also suggests the importance of quick feedback.

Koller et al. (2012) tracked hearer gaze using a camera and used that information to produce feedback to correct or confirm previous referring expressions. The experiment took place in a 3D virtual world. This enhancement was compared with feedback based on the instruction follower's position and with a system with no feedback at all. The eye-tracking enhancement significantly improved hearers' understanding of the REs. Eye-tracking is therefore a useful tool to improve interaction quality. This experiment also shows the importance of feedback, since the system with no feedback performed worse than the two systems with feedback.

Chapter 2

GIVE Challenge

An important framework for this thesis is the GIVE Challenge (Koller et al., 2010a). The data I used to study the referring strategies were collected using the GIVE framework. I used the GIVE framework to implement and test a hypothesis regarding these strategies. Therefore, in this chapter I will describe this academic competition in detail.

The first section will answer basic questions, such as what is the GIVE Challenge, why was it created and what are its interesting properties.

In the next section, I will provide a brief history of the GIVE Challenge together with some of its results.

In the third section, the focus will be a detailed description of the shared task and the virtual world of the GIVE Challenge.

2.1 Introduction

The GIVE Challenge was a series of Natural Language Generation (NLG) competitions run from November 2008 to March 2012. Participants developed NLG systems to navigate human-controlled avatars through a 3D virtual environment. The real-time navigation was realized through written instructions displayed on the screen. The goal of the navigation was to finish a treasure-hunt game. In Figure 2.1 we can see the GIVE client with a virtual world and an example of an instruction. A more detailed description of the task and the environment is in the Section 2.3.

Koller et al. (2010a) state that one of the goals of the GIVE Challenge was spawning an interest in NLG, a subfield of computational linguistics (CL), and was inspired by other competitions in the field such as the Recognizing Textual Entailment challenge¹ and NIST machine translation competition².

According to Koller et al. (2010a), another important goal was to introduce and

¹<http://pascallin.ecs.soton.ac.uk/Challenges/>

²<http://www.itl.nist.gov/iad/mig/tests/mt/>




Figure 2.1: A human subject is being navigated through the environment.

explore a new way of evaluating NLG algorithms, techniques and systems in a shared task. More specifically a shared task which was, on the one hand, complex enough to encompass multiple NLG subtasks and, on the other hand, was only concerned with NLG and not any other fields of computational linguistic.

Three basic approaches to evaluate NLG systems are to compare system output to annotated corpora, measuring task performance in an experiment and human judges evaluation. Koller et al. (2010a) argue the advantages and disadvantages of these evaluation methods in more depth, therefore I will only provide a brief summary. The first approach compares output of the NLG system to an annotated corpus, also known as a gold-standard. It is a fast and cheap approach, but a problem with it lies in the complexity of natural language. We can often express concepts in many different ways and there is often no telling which way is a better one. The second approach conducts an experiment and measures task performance on human subjects. Measuring task performance avoids the problems of the gold-standard, but it is expensive and time consuming. Lastly, trained human judges are used to evaluate the system. It is less demanding than the second approach, but for the cost of certainty, that the results correspond with results one might achieve with non-expert subjects.

The GIVE Challenge proposes and successfully implements a new approach, in a sense that it wasn't used for NLG before, through Internet-based evaluation. The basic premise is using a client-server software methodology. The client is a program installed on test subject computer, which is easily downloadable from a public website. The client connects through the Internet to a matchmaker server and a random evaluation world is selected. The matchmaker also connects the client to a randomly selected NLG system, which itself, can be hosted on a different server. Client and NLG systems then communicate back and forth until the task is finished.

The matchmaker finally logs the entire sessions to a database. Figure 2.2 shows that architecture in a simple diagram.




Figure 2.2: Software architecture of GIVE Challenge.

This approach immediately presents several advantages. It does not require the physical presence of the test subject in a laboratory. The subject simply downloads the client from a website and is able to do the experiment at his/her convenience. A second obvious advantage is a scalability. The number of individuals which can parallelly undertake the experiment is only limited by the servers' load. Thanks to the low costs, advertisement becomes the decisive limitation on the number of subjects. Take for example the second instalment of the GIVE Challenge which had up to 1800 participants.

On the other hand, part of the control over the experiment is lost in this approach; for example the control over the subject pool. Another problem which rises with this approach is that individuals can repeat the experiment.

In addition to the Internet-based evaluation, the GIVE Challenge utilized a variety of evaluation measures of both objective and subjective nature. Among the objective measures were task success rate, number of instructions and time required to finish the task. For subjective measures a questionnaire was used at the end of the session. The questionnaire mostly used a 5 point scale with questions such as how clear where the instruction or how friendly was the system. Some of the measures intentionally collided with each other, to motivate different approaches of generating instructions.

Having presented the basic concept of the GIVE Challenge and the reasons for its creation, I will now move onto a brief history of this competition.

2.2 History of GIVE Challenge

The first instalment of the GIVE Challenge (GIVE-1) was publicized in March 2008. Koller et al. (2010a) report on this instalment and are the source of the following information. For more details please refer to their paper. The data collection period was from November 2008 to February 2009. Four teams participated in this challenge, namely from these universities: University of Texas at Austin, Universidad

Complutense de Madrid, University of Twente and Union College. The team from University of Twente submitted two systems, making the final number of systems five.

What is important to note about GIVE-1 is a different world representation from the following instalments. GIVE-1 used a discrete square grid for player movement. The player was able to rotate only by 90° and walk forward and backwards by one square of the grid. That had a major impact on the design of the NLG systems. Participating teams at least occasionally used this grid in their references (eg. *move forward three steps*). Afterwards organizers realized that the grid and the discrete movement made the task easier than intended and they were after GIVE-1 removed.

Altogether, 1143 valid games were recorded. The demographics featured a majority of males (over 80%) and a wide spread over different countries in the world. For the actual results, the system from Austin significantly outperformed all other systems in task completion time. At the same time the systems from Union and Madrid outperformed the other systems in success rate. That shows the significance of different measures for the evaluation. Other interesting conclusions in both the objective and subjective measures can be found in previously mentioned paper. Apart from objective and subjective measures, the report examined the influence of English language proficiency and differences between the evaluation worlds. The English proficiency had an impact on the task success rate but solely for the least proficient category. The evaluation world also had a significant influence on the task success rate.

Finally, the first instalment also compared the Internet-based evaluation with a more standard laboratory evaluation. The conclusion was that Internet-based evaluation provides meaningful results comparable and even more precise in some areas to the laboratory setting.

The second instalment (GIVE-2) ran from August 2009 (data collection starting in February 2010) to May 2010 and is thoroughly described by Koller et al. (2010b). The following information is based on this paper. The biggest difference to GIVE-1, which was mentioned previously, is that players were now able to move freely. This made the instruction generation considerably harder. Additionally, the questionnaire was revised and a few new objective measures were introduced. The evaluation worlds used in GIVE-2 were considerably harder than in GIVE-1. The number of distracting buttons was increased and same-colored buttons were in some cases next to each other. Also the number of alarm tiles was increased. Otherwise, the architecture and the rest of the details stayed the same as in GIVE-1.

This time 1825 games were played over seven NLG systems developed by six teams from: Dublin Institute of Technology, Trinity College Dublin, Universidad Complutense de Madrid, University of Heidelberg, Saarland University and INRIA Grand-Est in Nancy (2 systems).

There was a big drop in success rate, most likely linked to the free movement and the increase of difficulty in the evaluation worlds. Similar to the results in GIVE-1, there was an influence of English proficiency and game world on the task success rate. Additionally, the age of the subjects played a role in the time required to finish the task and the number of actions to finish the task (younger subjects being faster

and requiring less actions).

Some teams participating in the GIVE Challenge tried to use a corpus of human to human interactions in the GIVE scenario. They were learning language expressions or decision-making processes and applying them in their NLG systems. The teams were however relying on small self-collected datasets. In a light of this, organizers of GIVE Challenge decided they would collect and provide dataset for future use. Gargett et al. (2010) describe this dataset, which was used in the next instalment of the GIVE challenge.

Following GIVE-2 was the so called Second Second instalment (GIVE-2.5), which kept almost the same settings as GIVE-2. There was just a small addition to the objective measures and a reduction in the number of subjective questions. The data collection took place between July 2011 and March 2012. Striegnitz et al. (2011) report on the partial results of 536 valid games from July and August 2011, which however constitute a majority of the final number of 650 valid games.

Eight NLG systems participated from 7 teams: University of Aberdeen, University of Bremen, Universidad Nacional de Córdoba, Universidad Nacional de Córdoba and LORIA/CNRS, LORIA/CNRS, University of Potsdam (2 systems) and University of Twente. In this instalment the teams employed a broader spectrum of approaches. The team from the University of Bremen used decision trees learned from the GIVE-2 corpus. Universidad Nacional de Córdoba and LORIA/CNRS, LORIA/CNRS selected instructions from a corpus of human to human interactions. The teams also often included algorithms from existing NLG and CL literature.

Apart from comparing the systems through objective and subjective measures, Striegnitz et al. (2011) again examined effects of evaluation worlds and demographics factors on task success rate. The evaluation worlds and the English proficiency had an effect. Additionally computer expertise and familiarity with computer games significantly influenced the task performance.

The following section describes the shared task in more detail and lists possible objects of the GIVE virtual worlds.

2.3 Task and GIVE world

The GIVE world is a 3D virtual world. The world is an indoor environment, comprising of rooms connected by doors. It's defined in a human-readable format and stored in a text file. The following objects can be places in a world:

- Alarm tile
- Button
- Door
- Landmark

- Bed
 - Chair
 - Couch
 - Dresser
 - Flower
 - Lamp
 - Table
 - Window
- Picture
 - Safe
 - Trophy
 - Wall

In addition, some of these objects can have attributes, states or can operate other objects. Buttons have colors as an example of an attribute. Doors and safes can be in a closed or an open state. Buttons can operate doors, safes or pictures.

Walls are actually created automatically by defining the shapes of rooms. Rooms can have a rectangular shape or can be defined by a polygon. I will sometimes use the term “corridor” which is a connecting room, usually not containing any button.

The landmarks serve as decoration but they can be used in referring expression generation. A picture is technically a landmark as well, but in the GIVE Challenge it often serves another purpose. It covers the safe and needs to be put aside by a button press.

Figure 2.3 shows an evaluation world from GIVE-2.5. In the top-left room we can see the player starting position. Buttons are colored squares on the walls. The grey bars on the walls are closed doors. The trophy is in a safe in the middle-left room. There are also landmarks (like lamp or chair) and one big red square marking an alarm tile.

The flexibility of the GIVE world creation allows a relatively broad range of scenarios for the task. On the other hand, all the GIVE Challenge instalments consisted of a similar sequence of steps. The goal of all the GIVE Challenge worlds is to pick up a trophy. The trophy is hidden in a closed safe. In order to open the safe a sequence of buttons, usually counting somewhere around 6 buttons, has to be pressed. The safe can be also hidden by a picture, which needs to be put aside. The buttons in a safe-opening series are often in different rooms. Rooms can be also closed off, requiring another button press to open the door. While moving around the world, the player has to avoid alarm tiles. Stepping on an activated alarm tile causes an immediate loss. Alarm tiles can block the only path and need to be deactivated by a button press. Some buttons also cause an alarm and an immediate loss.




Figure 2.3: Example GIVE world viewed in GIVE map viewer utility.

Depending on the number of rooms, the complexity of button arrangements and the length of the safe-opening sequence, the task can range from short and trivial problems easily handled by a few instruction templates, to a long and hard case, where it's impossible to capture every possible scenario.

To summarize, the navigation system of a player in the GIVE world has to deal with the following steps. Note that their order depends on the world definition and they can be thought of as layers of behaviours the system must enforce on the player.

- Avoid alarm tiles
- Avoid pressing alarm-causing buttons
- Deactivate path-blocking alarm tiles by a button press
- Open closed-off rooms by pressing the correct buttons
- Press a sequence of buttons to open the safe
- Reveal the safe behind a picture by pressing the correct button
- Take the trophy

After the safe was opened and possibly revealed from behind the picture the player can pick up the trophy and therefore win the game.

Chapter 3

The S-GIVE Dataset

After GIVE 2.5 instalment presented in Chapter 2, Prof. Kristina Striegnitz, Ph.D., one of the organizers of the GIVE Challenge, became interested in spoken communication and therefore decided to collect a new dataset, called the S-GIVE Dataset. This chapter serves as an introduction and analysis of this dataset.

As a side note, Striegnitz et al. (2012) report on a smaller German dataset, which is similar to the one I will be talking about.

In the first section, I will introduce the dataset and provide technical details of how it was created. Section 3.2, will, after the fashion of the GIVE Challenge look at how world and demographic factors influenced the task performance. The next section analyses REs in the dataset. The last section explores the phenomenon of chains of references.

3.1 General overview

The S-GIVE dataset is different from previous GIVE Challenge experiments because the IG's instructions were of a spoken form. That changes many aspects of the discourse. For one thing, the IG knows when the IF received his instruction, which is not true for the written instructions. That promotes faster feedback and allows interrupting during an instruction. However, in some cases, the spoken word tends to be less formal and grammatically correct. Moreover, interjections are very common, as are incomplete sentences. That makes the S-GIVE dataset complicated yet interesting to explore.

The data-collection started in July and finished in November of 2012. Through that period 21 interactions between two human subjects were recorded. Originally, 22 pairs participated, but one of the pair failed to finish the tasks and is excluded from the dataset. The subjects were asked to bring someone they know and they were financially compensated for the effort.

The set-up for the experiment is shown in Figure 3.1. One human subject was an instruction giver (IG). He/she is on the right in Figure 3.1. His/her role was

essentially the role of the NLG system in the GIVE Challenge. He/she was able to see a map of the world, which was updated in real-time and he/she got information about all necessary steps to finish the task. In addition, he/she was able to see the other person’s client screen. He/she communicated with the other person through a microphone and his/her goal was to navigate the other person through the world and make him/her finish the treasure-hunt.

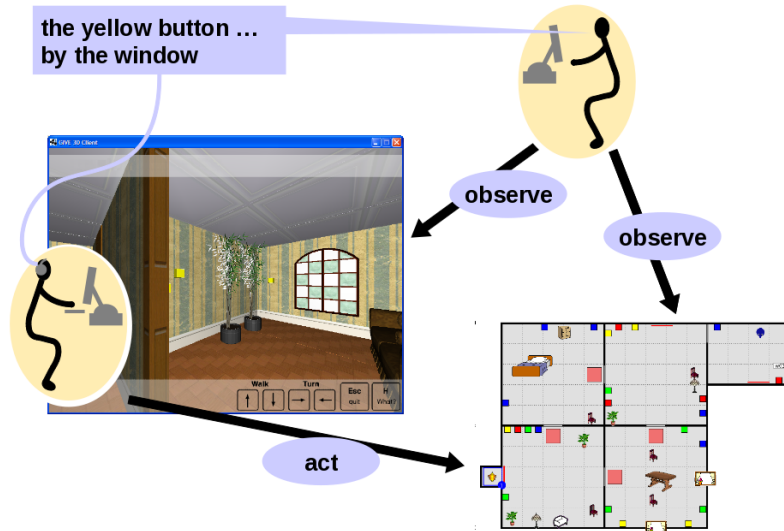


Figure 3.1: Experiment set-up of data collection

The other person was an instruction follower (IF). He/she is on the left in Figure 3.1. He/she interacted with the client and moved the avatar around the world and was able to press buttons. The IF listened to the IG’s instructions through a headset.

Each pair did one short tutorial world. After that they switched the roles of instruction giver and instruction follower. Following the tutorial was one “normal” world randomly chosen from two variants, marked world 1 and world 3 in the dataset. Maps of the worlds 1 and 3 are in Figures 3.2 and 3.3 respectively. Finally they did a difficult version of the other variant (if they started with world 1 the difficult version was for world 3 and vice versa). The difficult versions are marked 1-d and 3-d in the dataset. A difficult version of the world had an increased number of distracting buttons and landmarks compared to the “normal” version, as can be seen in the map of world 1-d in figure 3.4. If not present in the report or not stated otherwise, the short tutorial worlds are normally excluded from the analysis.

Similarly to the GIVE Challenge, after all 3 rounds subjects were asked to fill in a questionnaire. Its purpose was to get demographic and other relevant information on subjects. The questionnaire can be divided into three parts. The first part was only filled out by the IF and rated the IG and his instruction giving on a scale from 1 to 7. The complete list of the first part questions follows:

1. Overall, my partner gave me good instructions.
2. Interacting with my partner wasn’t annoying at all.




Figure 3.2: Map of the world 1 - normal version.




Figure 3.3: Map of the world 3 - normal version.




Figure 3.4: Map of the world 1 - difficult version.

3. My partner's instructions were clearly worded.
4. When I had problems with the instructions, we solved them quickly.
5. I enjoyed solving the task.
6. I felt I could trust my partner's instructions.
7. I really wanted to find that trophy.
8. My partner immediately offered help when I was in trouble.
9. I would recommend this experiment to a friend.
10. My partner's instructions were not repetitive.

The second part was filled by both the IF and the IG and was concerned with their navigation skills. To measure the ability to navigate in the world the Santa Barbara Sense of Direction (SBSOD) scores were used (Hegarty et al., 2002). Again the questions were on a scale from 1 to 7. Note that some of the question in the following list are of positive nature (higher rating equals better navigation skill) while other are of negative nature, therefore these scores had to be normalized.

1. I am very good at giving directions.
2. I have a poor memory for where I left things.
3. I am very good at judging distances.
4. My "sense of direction" is very good. I tend to think of my environment in terms of cardinal directions (N, S, E, W).
5. I very easily get lost in a new city.

6. I enjoy reading maps.
7. I have trouble understanding directions.
8. I am very good at reading maps.
9. I don't remember routes very well while riding as a passenger in a car.
10. I don't enjoy giving directions.
11. It's not important to me to know where I am.
12. I usually let someone else do the navigational planning for long trips.
13. I can usually remember a new route after I have travelled it only once.
14. I don't have a very good "mental map" of my environment.

Finally, the last part of the questionnaire was of a demographic character. We can see questions about age, gender, language and computer expertise, 3D games experience and knowledge of the partner in following list.

1. What is your age?
2. Are you male or female?
3. What is your profession / major / favorite subject in school?
4. How would you rate your computer expertise?
5. How familiar are you with 3D computer games?
6. How many hours per week do you play 3D computer games?
7. Was there a time in your life when you played more 3D computer games? If so, how many hours did you play then?
8. What languages do you speak? Please indicate how well you speak each on a scale from 1-5, where 5 is your native language.
9. Did you already know the second participant?
10. How well do you know the second participant?
11. Have you worked collaboratively with the second participant before? (For example, when doing homework or preparing a class presentation?)
12. Have you played 3D computer games with the second participant before?

Many of these questions are explored in the next section as potential factors influencing the task performance.

As was mentioned in Chapter 2, the entire session is logged to a database. The player’s position, orientation and all visible objects are logged at a fixed rate. Moreover, other information such as buttons presses or the end of the session are also stored in the logs. Because the worlds are static, distances and angles between the IF and other game objects are easily computed from these logs.

Apart from the logs, there are of course sound files of the IG giving directions. These were transcribed and together with some information from the logs transformed into ELAN files. ELAN is an annotation software (Sloetjes and Wittenburg, 2008). I will use the term *automatic annotations* for the ELAN files created from the logs and transcribed audio (however note that transcribing audio was done manually).

Building on top of these automatic annotations are manual annotations. They are primarily concerned with referring expressions and also stored in ELAN format. Most referring expressions in GIVE aim to locate a button, which needs to be pressed. I will call the button, which is a goal of a referring expression, the target button. Which button is the target button of a reference is one of the layers in the manual annotations. Another layer of the annotations is some basic categorization of the references, whether it is a reference to a single button, to a group of button, to a landmark and so on. The third layer looks deeper into the contents of the reference. It notes whether the reference contains for example the color of a button, whether distractors or landmarks are part of the reference or whether the reference explicitly points out that the button was already pressed before.

The previously mentioned logs, automatic annotations and manual annotations together form the dataset this chapter is dealing with.

An example of an interaction between an IF and an IG is in the following text. Spatial information is transcribed in parentheses for the sake of clearness.

(IF enters a room)
IG: Go towards the red buttons.
(IF turns right and start walking, but he turns too much)
IG: No the ones next to the lamp...
(IF corrects his direction)
IG: Yeah that lamp. On the right.
(IF is facing three buttons.)
IG: Press the button on the wall you are looking at, that’s far from the lamp and on the left.
(IF goes towards the correct button and stops close to him)
IG: Press it.

Next section, looks at how world selection and demographic factors affected the task performance.

3.2 World and Demographic factors

As was noted repeatedly in Chapter 2, the world had significant influence on the task success rate in GIVE Challenge. However in the GIVE Challenge the worlds were designed to be of a different difficulty, whereas in the S-GIVE dataset they were designed to be similar in terms of the difficulty to minimize effects outside of the navigation strategy. In addition, all the sessions were successful except for one which was discarded. The question about the influence of the worlds must be reformulated, since the task success rate no longer makes sense. Instead, I will measure task performance by time required to finish the task (duration) and also use other performance measure when appropriate.

Despite the design choices, I found that the normal worlds 1 and 3, had a different mean duration (p-value 0.0473 for unpaired two-sample t-test). There are multiple explanation for this difference. The relatively small number of subjects is certainly one of them. We can also notice in Figure 3.3 slightly complicated system of hallways in the center of world-3. But this discovery does not have a major influence on my work. Moreover, the difficult worlds did not have significant difference between their mean duration (p-value 0.6195 for unpaired two-sample t-test).

Another thing I was interested in was the influence of the gaming experience of both participants on certain performance measures, namely on the duration, on the average speed of IF movement and the time the IF spent moving. The average speed of the IF is simply a total distance the avatar controller by IF travelled in a session divided by the duration. The time the IF spent moving aggregates the time where avatar wasn't either motionless or wasn't only rotating in place.

I found correlations between gaming experience and these variables. Table 3.1 shows the correlation matrix between gaming experience and performance measure. Not surprisingly, these correlation are especially high for the IF, since he/she is the one who is actually playing the world, but the IG gaming experience seems to have some effect as well. The past gaming experience (questions 7 in third part of the questionnaire) is more important than contemporary playing (question 6 in the third part of the questionnaire). Most prominent are the familiarity of IF with 3D games (question 5), the IF hours per week spent playing games at the past peak gaming period (question 7), the same variable for IG and hours spent gaming per week for IF at present (question 6). For the difficult worlds some correlations change slightly. IF's gaming at past peak period has much weaker correlation with duration here. In general, individuals who are familiar with games (gamers) take less time to finish the world, they spent more time moving and they have higher speed.

VERIFY CORRELATIONS

I also looked at the influence of SBSOD scores (second part of the questionnaire) on the task proficiency. A correlation matrix revealed weak or almost no correlation between SBSOD scores and the time needed to complete the world, as can be seen in Table 3.2. In the difficult worlds, however, the correlation became slightly stronger.

NOTE ADD T-TEST

World	Question	Duration	Speed	Time moving
Normal	IF hours/week peak gaming (7)	-0.411	0.486	0.428
	IF hours/week now (6)	-0.366	0.338	0.255
	IF familiarity 3D games (5)	-0.590	0.720	0.664
	IG hours/week peak gaming (7)	-0.359	0.450	0.435
	IG hours/week now (6)	0.079	0.155	0.180
	IG familiarity 3D games (5)	-0.230	0.221	0.224
Difficult	IF hours/week peak gaming (7)	0.111	0.199	0.135
	IF hours/week now (6)	-0.388	0.312	0.233
	IF familiarity 3D games (5)	-0.478	0.569	0.520
	IG hours/week peak gaming (7)	-0.287	0.715	0.715
	IG hours/week now (6)	0.149	0.348	0.420
	IG familiarity 3D games (5)	0.009	0.403	0.473

Table 3.1: Correlation matrix of gaming experiences and performance measures

World	Role	Duration
Normal	IF	-0.085
	IG	-0.082
Difficult	IF	0.162
	IG	-0.210

Table 3.2: Correlation between SBSOD scores and duration

The data suggest that there is positive correlation between male gender and task proficiency measured as duration. Table 3.3 shows these correlations. The effect of male IG diminishes in the difficult worlds but the effect of IF is even stronger in the difficult worlds. However there are several facts to take in consideration here. First of all, we don't have enough data to have a statistically significant conclusion. This correlation might have also been caused by having more male gamers than female gamers. In fact, Table 3.4 suggest so. Lastly, there has been research about influence of gender on spatial cognition and mental rotation; an example of more recent one is (Geary et al., 2000). They conclude that males are more proficient in tasks requiring mental rotation. Since the IGs have to do mental rotation while giving direction in GIVE scenario, this might be a source of the correlation. Another paper worth considering on this topic is (Moffat et al., 1998), which found a gender difference in time required to finish a virtual maze.

World	Role	Duration
Normal	IF	-0.234
	IG	-0.277
Difficult	IF	-0.349
	IG	-0.106

Table 3.3: Correlation between male gender and duration

Role	Familiarity with 3D games
IF	0.341
IG	0.661

Table 3.4: Correlation between male gender and familiarity with 3D games

Table 3.5 shows correlation matrix for age. The age of the IF has a positive correlation with task proficiency measured in duration, and in difficult worlds this correlation is one of the strongest ones. Older IF are also moving less and are generally slower. For IG the correlations have the same direction, however they are much weaker.

World	Role	Duration	Speed	Time moving
Normal	IF	0.175	-0.467	-0.448
	IG	0.275	-0.098	-0.125
Difficult	IF	0.614	-0.275	-0.196
	IG	0.016	-0.217	-0.232

Table 3.5: Correlation matrix of age and performance measures

ADD CORRELATIONS OF AGE WITH GAMING EXPERIENCE

Lastly I was interested how familiarity of participants with each other (questions 9-12 in the third part) influenced the task performance. Table 3.6 shows, that knowing

the partner had a positive impact on task efficiency. The most correlated question was question 10: how well they know each other.

World	Question	Duration
Normal	IF Co-players in past (12)	-0.038
	IF Collaborative work (11)	0.109
	IF how well know each other (10)	0.529
	IF know each other (9)	0.164
	IG Co-players in past (12)	-0.124
	IG Collaborative work (11)	0.189
	IG how well know each other (10)	0.420
	IG know each other (9)	0.177
Difficult	IF Co-players in past (12)	-0.078
	IF Collaborative work (11)	0.170
	IF how well know each other (10)	0.437
	IF know each other (9)	0.247
	IG Co-players in past (12)	0.291
	IG Collaborative work (11)	0.245
	IG how well know each other (10)	0.361
	IG know each other (9)	0.189

Table 3.6: Correlation between participants familiarity and duration

Exploring interesting correlations, I will now focus on RE.

3.3 Referring expressions

Because REs are the main focus of my research, this section serves as a brief introductory analysis of REs in the S-GIVE dataset.

Overall, 793 REs were annotated in the manual annotations. Apart from the time interval of the reference, several other facts were annotated in the manual annotations, as was mentioned in Section 3.1. First of all, the target button of each RE was annotated. The count of distinct target buttons is 29.

REs were also separated into 5 high-level categories depending on what is the target of the reference. The overview of the categories is in the following list:

- Target - Referring to the target button
- Group - Referring to a group of buttons, one of which is the target button
- Landmark object - Referring to a landmark (any object or room feature, but not a button) which will then be used to locate the target button
- Landmark button - Referring to a distractor button as a landmark

- Remove button - Referring to a distractor button to exclude it

The percentage count of the categories is given in Table 3.7. References to target buttons are a dominant category. Around 10% of the references are group references. References to landmarks occupy only 6% of all references.

Category	Percentage (%)
Target	82.47
Group	10.34
Landmark object	5.04
Landmark button	1.39
Remove button	0.76

Table 3.7: Percentage of REs in the categories

Another layer of manual annotations looked into the content of the REs. It was done through annotating several semantic elements of the REs of the Target category. They are listed in the following list and their usage is summarised in Table 3.8. Please note that these elements are not necessarily exclusive with each other.

- Type - RE expressed the target object as its type (“button”)
- One - RE expressed the target object as “one”
- Pronoun - RE expressed the target object as a pronoun
- Color - RE contained color of the target object
- Button location - RE contained relative location of the target object to a distracting button
- Object location - RE contained relative location of the target object to a distracting object (not a button)
- IF location - RE contained relative location of the target object to the IF
- Room location - RE contained relative location of the target object to a room
- History - RE informed that the target object was already manipulated with

From previous tables, we can see prevalence of relatively simple RE. However, the next section will show that the situation is more complex than it may seem from this section.

Element	Percentage (%)
Type	55.35
Color	47.09
One	26.29
Button location	18.80
Object location	16.67
IF location	11.62
Pronoun	10.70
History	6.88
Room location	5.81

Table 3.8: Percentage of REs which contained a semantic element

3.4 Chains of references

An interesting phenomenon I have noticed and further explored in the dataset are consecutive references to one button. It can be seen in the following sentences: “Straight ahead of you there on the opposite wall there are two blue buttons. Press the one on the right. The one close to the picture.” The IG started off with a references to a group of two buttons; the target button being one of those two. In the second sentence he picked out the target button from the group. In the last sentence the IG made another reference containing a landmark, adding redundant information. Since the references are concerned with one target button and follow each other relatively fast, I have called them chains of references (chains, in short). These chains are examples of non-standard referencing strategies, I was mentioning since the beginning of this paper.

The chains vary in length, from short ones, consisting of only two references, up to lengthy ones with eight references following each other. The example from previous paragraph is three references long. Figure 3.5 shows a histogram of the chains’ length.

The chains are significant part of the discourse, in fact over 77% of all RE belong to a chain. However, the chains have multiple linguistic functions, which makes them difficult to explore. Even inside one chain, there are often combination of functions. I have manually annotated some common functions of the chains. The most common one is to inform the IF that he is supposed to press the target button. A simple example of this function, which I call action function, is in following discourse: “The red button in front of you. Press that one.” It may sound redundant to use action function as the experiment progresses, since the IF are not manipulating the buttons in any other way than pressing them. But I have found out that it’s often the case, that IG still use action functions in later stages of navigation.

Another common function is confirmation: “That same red button we pressed before, we’ll press that again. Yeah that one.” The IG confirms at the end of the previous utterance that IF is looking at or heading to a correct target button.

The IG often utters a RE, which does not perfectly “pick out” the target button from the set of buttons in the room. The IG can make up for that information




Figure 3.5: Histogram of the chain length

deficit with a confirmation or by further specifying with another RE. That is a specification function, as in following two sentences: “Now that green button. You want the one closer to the lamp.”

When a group of buttons is utilized in a RE and the target button is part of that group, it is inevitable that the IG will have to make another RE to “pick out” the target button from the group. Therefore group references imply chains of references and should be considered as one of the functions. A simple example of the group function: “Two blue buttons on the wall. Hit the blue button on the right side.”

When the IF has clearly chosen the wrong button, the IG will try to correct that error. I call that an error function. The following extract features this function: “That blue button. No no no. The other. That one.”

A summary of the functions, which I have manually annotated, is in Table 3.9. Please note that these functions are not exclusive. One chain can have both a confirmation and a specification function.

Function	Chains containing it (%)
Action	66.16
Confirmation	29.29
Specification	25.25
Group	24.74
Error	11.11

Table 3.9: Percentage of chains having specific functions

These chains will be further explored in the following chapter, but as one can already notice, they are complex language phenomenon.

Chapter 4

Machine Learning on REs

REs are big part of language generation of a navigation system and especially so in the GIVE scenario. My work attempts to apply machine learning techniques to the S-GIVE dataset with the goal to help the navigation system with REG. These attempts are presented in this chapter.

Machine learning (ML) is becoming a popular topic together with a very broad term of Big Data. This branch of artificial intelligence encompasses various problems such as classification or clustering. Research in the area of ML spawned various algorithms for solving these problems and evaluation techniques for comparing the models outputted by the learning process.

It is important to emphasize that ML creates models which can only be as good as the data they were trained on. Careful and thorough data analysis is crucial part of the ML process. When it comes to feature extraction, experimentation are sometimes the only methods for finding better models.

In the first section I will describe attempts at predicting the timing of the first reference to a target button.

The second section talks about modeling chains of references.

The third section briefly touches on the topic of using room memory.

In the last section, I will present my thoughts on the results of the previous sections.

4.1 Timing of the first reference

The work of Stoia et al. (2006a) was previously mentioned in the related work section of Chapter 1. They applied machine learning to the timing of the first reference in a 3D virtual world. The set-up of their experiment is quite similar to the GIVE scenario and so I decided it would be interesting to replicate their methodology on the GIVE dataset.

I defined the problem of the timing of the first reference as a classification task, as did Stoia et al. (2006a). More precisely binary classification, the two classes being

either refer to the target button or delay the reference. After extracting the first references to buttons which needed to be pressed from the dataset, I excluded plural references, because of their complexity. Some buttons were placed on top of each other and the IG wasn't sure which one need to be pressed. These were excluded as well because of the unnecessary confusion. In addition, some buttons were referred to multiple times, because they needed to be pressed twice. The later references may have a completely different structure and the buttons were therefore excluded. That left me with 351 first references. For each first reference I have chosen one negative example, where the IG could refer to the target but chose not to. I picked negative examples randomly from the interval between entering the room and the time of the first reference. Overall, that is 702 data-points with perfectly balanced classes.

As for features extraction, I have chosen similarly to Stoia et al. (2006a) various spatial information. For the positive examples, I averaged these spatial information over a 0.6 seconds window centered on the time of the reference. The reasoning for that, is that the IG takes the scene situation into consideration before and possibly after they start uttering the reference. For negative examples I chose not to average them, since they are chosen randomly. All features are listed in the following list. The list also includes figure numbers. These figures are histograms of the attributes, separated for both classes and can be found in appendix.

- Distance to target button - Figure 5
- Absolute value of angle to target - Figure 6
- Whether target is visible (True/False) - Figure 7
- Number of distractors - Figure 8
- Number of distracting buttons - Figure 9
- Number of visible buttons with smaller angle to IF than the target button - Figure 10

Once I have extracted these features I used three machine learning techniques: C4.5 decision trees because of their easy interpretation, naive Bayes to observe the effect of all attributes and Support vector machine for linear classification. I used the Weka software implementation of previous algorithms (Hall et al., 2009).

For all the algorithms I used a standard ten-fold cross validation. The results can be seen in table 4.1. A pruned decision tree for timing of the first reference can be seen in Figure 4.1. I used two simple baselines to compare the results with. I have perfectly balanced classes so the first baseline is 50%. A simple rule for the first reference is to refer when the target is visible, and delay it if the target is not visible. In my case that rule has an accuracy of 64.2%.

Only one algorithm was able to get over the visibility baseline and not by a significant amount. These results were surprising, because Stoia et al. (2006a) had success with the same approach on a similar dataset. Reasons for this difference are probably in the differences between their experiment and the GIVE set-up. First, their tasks also

Model	Accuracy (%)
Class baseline	50.00
Visibility baseline	64.2
Naive Bayes	64.70
C4.5	63.31
SVM	55.42

Table 4.1: Results of first reference timing modeling

included different actions than pushing buttons (e.g. picking up items). Second, their worlds had higher diversity of the distractors and smaller frequency of them. With increasing number of distractors and particularly distractors of the same category, it seems that spatial features loose their power in predicting the timing of the first reference. The number of visible distractors was the best attribute in their decision tree, but in my tree (Figure 4.1) it had lower information gain. Moreover, the decision tree did the split on the number of visible buttons not the number of all visible distractors.

After the timing of the first references classification proved to be more difficult than expected, I have switched from predicting the timing of the references to predicting their content. Next, I focused on chains of references.

```

Distance to target <= 4.45
| Is target visible? = False: DELAY
| Is target visible? = True
| | Distance to target <= 2.51: REFER
| | Distance to target > 2.51
| | | Number of distracting buttons <= 4.67: REFER
| | | Number of distracting buttons > 4.67: DELAY
Distance to target > 4.45: DELAY

```

Figure 4.1: Decision tree for first reference timing

4.2 Chains of references

Section 3.4 introduced the phenomenon of chains of references. It also analysed various linguistic functions the chains can play in REG. This section will build on top of this analysis, by employing machine learning techniques to model the chains.

A valid and important question is whether chains aren't something, which should actually be avoided, instead of modelled. That is, what is the relation between usage of chains and task performance measure, such as duration. To address this issues, I used linear regression predicting the duration of the experiment explained by the average chain length. Figure 4.2 does contain a hint of a trend, but also contains a lot of counter-examples.




Figure 4.2: Scatter plot of duration and chain length

When I applied linear regression the R^2 was 0.188, which means the average chain length explains very little of the duration variation. If the average chain length was to significantly influence the duration, I would expect much higher R^2 . The correlation between the duration and average chain length is not low: 0.433, however correlation does not imply causation. Longer chains can be caused by errors of the IF or the IG and the errors are also likely to increase the duration. Based on these facts, I don't believe the chains are harmful phenomenon and it makes sense to proceed with attempting to model them.

With modeling the chains of references, there are numerous problems to tackle. The simple presence of a chain can be thought of as a classification task. It can be interesting to try to predict the chain length. Having established common linguistic functions of the chains, classifying chains whether they would contain these functions is another task to consider. I looked into each of these problem. Trough extracting relevant features I thoroughly explored using machine learning for modeling the chains.

As a side note, while exploring chains of references I switched from the Weka implementation of machine learning techniques (Hall et al., 2009) to scikit learn package for Python (Pedregosa et al., 2011). The main reason being easier automation of the entire pipeline, therefore speeding-up of the whole process and also more control over the whole process.

4.2.1 Presence of the chains

From Section 3.4, the functions the chains can play in a discourse are known. However, what role does the scene complexity and particular game state play in the presence of the chains? Will more complex scenes spawn more chains or are the chains too complex to be captured by simply looking at where they are created? To answer these question, I extracted scene complexity features and the target button for all references and classified whether a chain was present or not.

The features I extracted from the dataset are in the following list:

- Target button
- Number of objects in the room - Figure 13
- Number of buttons in the room - Figure 14
- Number of landmarks in the room - Figure 15
- Number of very close buttons to the target button (closer than 0.3m) - Figure 16
- Number of buttons on the same wall as target button - Figure 17
- Number of close buttons to the target button (closer than 1m) - Figure 18
- Number of far buttons (farther than 1.5m) - Figure 19

The target button is a categorical feature, so I used DictVectorizer class in scikit to transform it into multiple numerical features. I further tried to use all features, to select the 8 best features using the SelectKBest class from scikit with the χ^2 statistic as a scoring function and to select the 16 best features using the same strategy. As algorithms I used Decision Trees, Naive Bayes, Support Vector Machines (SVM), One-Nearest Neighbour (1-NN), Two-Nearest Neighbours (2-NN) and Random Forest. This selection employs various approaches to classification tasks, each of the algorithms having strong and weak points. For evaluation, I used standard ten-fold cross validation and compared the mean accuracy of the classifiers with the majority class baseline. From now on, I will also include double the standard deviation (std). The idea behind the double is that if the features follow a normal distribution and the three sigma rule is applied, the chance of the accuracy being in interval of \pm standard deviation times two is 95.45% . The results are summarized in Table 4.2.

Features	Model	Mean accuracy (%)	2 \times std
Majority baseline		55.2	
All	Decision Tree	55.5	13.5
	Naive Bayes	55.8	6.6
	SVM	56.1	11.2
	1-NN	56.3	9.3
	2-NN	53.6	12.1
	Random Forest	55.0	13.6
8 best	Decision Tree	59.7	14.6
	Naive Bayes	56.4	2.9
	SVM	58.3	13.1
	1-NN	58.8	11.1
	2-NN	56.1	14.4
	Random Forest	60.0	16.6
16 best	Decision Tree	55.8	13.5
	Naive Bayes	56.4	2.9
	SVM	58.3	13.9
	1-NN	56.9	9.1
	2-NN	53.3	12.8
	Random Forest	57.2	12.8

Table 4.2: Results of chains presence modeling

The best mean accuracy had the Random Forest for 8 best features. However, it wasn't significantly better than the majority class baseline. Taking into account the standard deviation, none of the classifier and features combination outperformed the baseline. From these results, I conclude that the chains presence is not dependent only on the scene complexity and specific scenarios. IG strategies for referencing, IF behaviour and other circumstances play role in the creation of the chains.

4.2.2 Chain length prediction

After not being able to predict presence of the chains based on spatial information and the target button, I was interested if I could predict the chain length. I used the same attributes as in the previous classification of the chains' presence, but added two more features concerned with IF movement behaviour. Namely, the ratio of the time the IF spent not moving at all through the chain duration and ratio of the time IF spent only rotating in place through the chain duration. The idea behind these features is that an IF who is not moving or just looking around is an indication of him/her being confused, which then should produce more referring expression for the chain.

For reference, I list all features and add figure numbers for the two new attributes:

- Target button
- Number of objects In the room
- Number of buttons in the room
- Number of landmarks in the room
- Number of very close buttons to the target button (closer than 0.3m)
- Number of buttons on the same wall as target button
- Number of close buttons to the target button (closer than 1m)
- Number of far buttons (farther than 1.5m)
- Ratio of time IF spent not moving - Figure 11
- Ratio of time IF spent rotating - Figure 12

I applied linear regression, predicting chain length based on the three groups of attributes mentioned - the target, spatial features and IF movement behaviour features. I evaluated the regressions by looking at R^2 . The results are in Table 4.3.

Features	R^2
Target button	0.154
Spatial	0.146
IF movement	0.126
Spatial ad IF movement	0.262

Table 4.3: Results of chains length modeling

None of the regressions from Table 4.3 are particularly good at predicting the chain length. Combining spatial and IF movement features does increase the percentage of chain length variation explained by the model, suggesting that these features have, an however small, effect on the chain length. Once again it shows that chains are complex phenomenon, influenced by IF, IG and scene variables.

4.2.3 Closer look at chains' content

Despite not being able to predict the chains presence and length, I was still interested in chains and decided to look closely into the chains content. First, I took advantage of having annotated the functions in the chains, as introduced in Section 3.4. I focused on the specification and group function and tried to predict whether the chain will contain these functions. Second, I tried to predict whether the chain provides new information about the target button after its first reference. For example the IG could add a RE about the position of the target button relative to a landmark in the third reference of the chain. The reasoning behind this classification is trying to predict when the IG adds information to the chains.

For all these classification tasks, I used the same features as in the chain length prediction, that is: the target button, spatial features, IF movement behaviour features and a combination of the previous two. Again, I used ten-fold cross validation for evaluation and compared that to the majority class baseline. The results for the specification function are in Table 4.4, for the group function in Table 4.5 and for predicting new information in Table 4.6. As algorithms, I maintained a broad range of algorithms, similarly to the previous machine learning attempts.

Features	Model	Mean accuracy (%)	2× std (%)
Majority baseline		74.7	
Target button	Decision Tree	73.3	6.0
	Naive Bayes	38.0	22.4
	SVM	74.7	5.3
	1-NN	70.7	14.8
	Random Forest	74.0	7.2
IF movement	Decision Tree	64.0	21.7
	Naive Bayes	72.7	11.1
	SVM	74.7	5.3
	1-NN	77.3	12.2
	Random Forest	71.3	16.9
Spatial	Decision Tree	74.7	16.7
	Naive Bayes	66.7	20.7
	SVM	71.3	14.7
	1-NN	68.7	23.9
	Random Forest	74.7	18.7
Spatial and movement	Decision Tree	66.7	23.1
	Naive Bayes	66.7	20.7
	SVM	71.3	18.9
	1-NN	76.7	8.9
	Random Forest	68.0	16.7

Table 4.4: Results of specification function in chains modeling

The best classifiers for these 3 tasks can be found in the respective results tables,

Features	Model	Mean accuracy (%)	2× std (%)
Majority baseline		75.4	
Target button	Decision Tree	81.3	14.4
	Naive Bayes	38.0	14.7
	SVM	75.3	6.1
	1-NN	79.3	13.9
	Random Forest	79.3	17.3
IF movement	Decision Tree	61.3	23.7
	Naive Bayes	75.3	6.1
	SVM	75.3	6.1
	1-NN	73.3	15.8
	Random Forest	68.0	16.7
Spatial	Decision Tree	80.0	10.3
	Naive Bayes	77.3	18.1
	SVM	77.3	8.8
	1-NN	77.3	13.6
	Random Forest	82.0	13.4
Spatial and movement	Decision Tree	71.3	16.9
	Naive Bayes	78.0	15.8
	SVM	77.3	8.8
	1-NN	79.3	9.3
	Random Forest	74.0	18.3

Table 4.5: Results of group function in chains modeling

Features	Model	Mean accuracy (%)	2× std (%)
Majority baseline		67.3	
Target button	Decision Tree	68.3	19.6
	Naive Bayes	38.0	15.8
	SVM	67.3	4.0
	1-NN	56.0	31.1
	Random Forest	68.0	19.6
IF movement	Decision Tree	58.7	8.0
	Naive Bayes	68.0	11.6
	SVM	67.3	4.0
	1-NN	41.3	20.5
	Random Forest	57.3	12.2
Spatial	Decision Tree	68.0	16.7
	Naive Bayes	60.7	28.3
	SVM	71.3	8.5
	1-NN	59.3	21.9
	Random Forest	67.3	17.3
Spatial and movement	Decision Tree	66.0	26.3
	Naive Bayes	62.7	26.1
	SVM	72.7	9.3
	1-NN	59.3	19.3
	Random Forest	64.7	18.9

Table 4.6: Results of new information in chains modeling

marked by bold font. Similar to the previous attempts, all combinations of attributes and algorithms had an accuracy around the baseline. The last section of this chapter will address these and all previous results in more depth. But before that, I modeled one more subtask, which is introduced in the following section

4.3 Room memory

The last machine learning problem I tackled was using memory of the previously visited rooms in the RE for room switching. When the IF has to return to a recently visited room, the IG sometimes employs REs such as: “Go to the room you were just in.” Deciding whether to use memory in room switching RE can be thought of as a classification task.

I extracted all moments where the IF went between rooms and determined whether the room memory was used from what the IG was saying before going into the new room. I also extracted 3 features:

- Was IF in the new room before?
- How many seconds before, was he/she there?
- How many rooms before, was he/she there?

The last feature being of a categorical character, once again DictVectorizer was used for transformation and from the transformed numerical values I selected the 3 best with SelectKBest and the χ^2 statistic. The results are in Table 4.7.

Model	Mean accuracy (%)	2× std (%)
Majority baseline	78.5	
Decision Tree	78.0	7.1
Naive Bayes	68.3	7.3
SVM	82.2	6.2
1-NN	80.8	6.6
Random Forest	84.4	6.6

Table 4.7: Results of room memory modeling

Despite not being able to significantly outperform the baseline using machine learning, I used the knowledge gained in the room memory modeling attempt. In the systems I developed for the experiment, which will be discussed in the last chapter, I created a simple rule for using room memory and enhanced the REG process.

The next section analyses results from previous sections and provides my thoughts on them.

4.4 Thoughts on ML in S-GIVE dataset

Thanks to better availability of annotated corpora, machine learning techniques are becoming popular in the field of REG, examples being mentioned in Section 1.2. Inspired by similar research, I looked into a S-GIVE dataset, so far unexplored using ML, a dataset of spoken instruction giving in 3D virtual environment.

Despite extensive feature extraction and tackling various problems, the results were unsatisfactory. In this section, I would like to present my thoughts, why I believe that was the case.

First point to explore is the complexity of the S-GIVE scenario. As was already discussed in Section 4.1, the S-GIVE worlds have specific properties, which most likely have significant influence on the way the IGs formulate their instructions. The S-GIVE dataset is specific in the high frequency of distractors, especially of the same category as the target of the references. To identify targets in such an environment often requires taking into account relations between entities. This was one of the original constraints of early REG research, as mentioned in Section 1.1, and only recently this constraint was attempted to be lifted. In some scenarios, the buttons were intentionally organized in complex arrays, where even the IG wasn't sure which button needs to be pressed. Relatively simple spatial features, which proved to be effective in research of Stoia et al. (2006a), likely lose their predictive power when dealing with situations of frequent distractors of the same type as the target.

The dataset size has to be taken into account, when discussing the ML results. The relatively small number of participants and, on the other hand, higher number of scenarios, where REs were created, is something to consider. With a smaller number of participants, the IGs' personal strategies of referring will play a bigger role in the dataset. Some might prefer to refer to the target button immediately, while others might postpone the reference until they removed more distractors through other movement instructions. This difference obviously makes any attempts of machine learning difficult on a small dataset. With enough data it would be possible to separate strategies to clusters and examine each cluster separately. However with 20 participants, this is not possible or very difficult to achieve. These strategies are also likely to be influenced by the IFs, further increasing the uniqueness of the REs produced for each pair of IG and IF. I strongly believe a bigger dataset would be more successful in applying ML techniques.

Last but not least, there is a question of the features used. Spatial features are easier to extract from S-GIVE dataset than any other. While I attempted using other features apart from the spatial ones, their extraction is problematic and time consuming. I would speculate that features about the IF's behaviour and the IG's personality might increase accuracy of presented models, but exploring them was beyond the scope of this thesis. Not to mention that some interesting features of a more personal character don't necessarily have to be in the dataset at all and one would have to repeat the experiment to collect them.

That being said, the ML attempts inspired the systems I used in the experiment described in the following chapter in many ways.

Chapter 5

Experiment on RE

Throughout this thesis, I was interested in a strategy of referring, where the first reference does not always uniquely identify the target object. Instead, the strategy relied on a feedback and additional references, which together with the first reference formed a chain of references. I wondered what is the effect of this splitting of the information across more references. In Chapter 4 this strategy was explored through phenomenon of chains of references. After not being able to model this behaviour through methods of machine learning, I decided to look at that strategy under the more controlled conditions of an experiment.

In the first section, I'll present my hypothesis concerning the strategy of splitting references.

The second section will describe the experimental set-up.

In the last section, I'll summarize the results of the experiment.

5.1 Hypothesis

To evaluate the strategy of splitting references, I compared two NLG systems, which I have built to resemble human IGs from the S-GIVE dataset. The two systems are identical, except for a difference in the REG strategy for the target button. One system represents the strategy of splitting references, while the other one would represent a more standard approach of referencing. They will be thoroughly described in Section 5.2.

A question this chapter tries to answer is how splitting references affects the task proficiency. Is this strategy significantly different from the more standard fully-identifying references? Or is it simply a language device the IG can utilize in complicated virtual worlds.

I've formulated my prediction to these questions as hypothesis:

Hypothesis 1. *Distributing information to uniquely identify the referent across multiple referring expressions separated by non-negligible time intervals does not*

have an impact on the task proficiency.

Measuring task proficiency can be done in several ways. In the previous chapters, I've often employed the duration of the experiment. However, for this experiment I've chosen a more specific measure, in order to diminish effects of other factors. I've measured the time from the point of uttering the first reference up to the pressing of the targeted button. This not only diminishes effects of other variables in the navigation process, but can also be interpreted as a measure how well the references were understood by the IF.

For testing the hypothesis I've compared average times between the first reference and the pressing of the target button for both systems. I have used the independent two-sample t-test for unequal sample size. The unequal sample size comes from the fact that a subject could press a wrong button and therefore the number of potential references may vary.

Having formulated hypothesis and translating it into a statistical test, I can now describe how I approached setting up the experiment.

5.2 Experimental set-up

I've tested my hypothesis through a human subjects evaluation. Each subject did 5 virtual worlds from the GIVE scenario. One of the worlds was a short tutorial world. This tutorial world was excluded from the analysis and served the purpose of diminishing the effect of learning. After the tutorial world, the subjects did 4 evaluation worlds in a random order. The two tested NLG systems were assigned to each world semi-randomly to ensure both of the systems appeared twice. Please note, that the tutorial world was created in such way that the two tested NLG systems behaved the same way and therefore none of the systems had an advantage in number of trials. I've used text instruction presented on the screen as in the GIVE Challenge instalments. Using some sort of speech synthesis or even recorded instructions was beyond the scope of this thesis.

The worlds were similar to the GIVE challenge worlds. Maps generated by the GIVE map-viewer for all of the 4 evaluation worlds can be found in Figures 5.1, 5.2, 5.3 and 5.4. I did not include any alarm tiles or alarm-causing buttons in the worlds, to avoid situations where the IF would loose. I also avoided complex arrays of buttons, which were present in some of the S-GIVE worlds.

I have called the two NLG systems Alpha and Beta. Since they are almost identical except for a REG strategy for the buttons, I will describe the Alpha system and simply point out the differences between it and the Beta system, at the end of the description. As was previously mentioned in Chapter 4, the NLG systems are inspired by the analysis of the S-GIVE dataset. Even though the S-GIVE dataset is spoken, I managed to transform some ideas from the spoken data to the written instructions.

Because the final systems are not the subjects of this thesis, but merely a device




Figure 5.1: Map of the evaluation world 1.




Figure 5.2: Map of the evaluation world 2.




Figure 5.3: Map of the evaluation world 3.




Figure 5.4: Map of the evaluation world 4.

to examine a NLP phenomenon, I will stick to a high-level descriptions of the ideas behind them and avoid exhaustive software engineering description, such as class diagrams and similar devices.

First thing I would like to address is the timing. In the spoken S-GIVE scenario the IG knows when the IF heard the RE. That is however unknown for the written instruction giving. I have partially solved this problem by estimating the timing with an average reading speed of 150 words per minute. This can of course affect the task proficiency because of personal differences in the reading speed and effects of learning, but it is not easily solved completely in the framework I was using.

For the direction giving, the Alpha system uses 4 directions. They were heavily used in the S-GIVE dataset even though sometimes enhanced by adjectives. The directions are simple: in front of the IF, left, right and behind the IF; all from the IF's point of view. In conjunction with REs which use relations between the world's entities, these 4 directions are sufficient to describe the path for the IF. Suppose we consider 0° a direction the IF is facing, then I've chosen the limits for the 4 directions as follows: in front $\langle -35^\circ, 35^\circ \rangle$, right $\langle 35^\circ, 145^\circ \rangle$, behind $\langle 145^\circ, -145^\circ \rangle$ and left $\langle -145^\circ, -35^\circ \rangle$. The system uses them as in the following example: "Press the green button on your left."

Before focusing on RE to the target button, I will report on how I implemented navigating through the rooms. In relatively simple scenarios with rooms and mainly straight corridors, the system can simply create a RE for the door leading to the next planned room, once the IF entered a new room. Adding verbs of movement such as "go" to this RE is a simple, yet in the S-GIVE dataset common method for navigating in GIVE scenario. An example of such a sentence is: "Go through the door in front of you." The only complication is that with only 4 directions, other doors can be present in the same direction as the target door. I've solved this by another RE, identifying the target door using its relative position in the group of distracting doors: "The door closest to you." Whenever this specification is needed, I've also added positive feedback when the IF is heading towards the correct doors and negative feedback if he/she enters a wrong room. I have also enhanced this subsystem by two additional improvements commonly seen in the S-GIVE dataset.

First, when the IF is only passing through the room on his way to a next one, I modified the language realization to take advantage of that fact. The system will produce expressions such as: "Make a left." or "Keep going straight." This adds variability to the NLG system and makes it more human-like.

Second, I exploited the room memory, first presented in Section 4.3. When the IF is immediately returning to the room he/she was just in and the time elapsed since he/she was there is not large (less than 20 seconds), the system will produce an expression such as: "Go back to the room you were just in." This addition not only increase human-likeness, but I would argue also has an impact on the task efficiency, since it avoids a lot of problems in the standard method described above.

The most complicated task are REs to a target button. This is where the systems Alpha and Beta differ. Once the IF enters the room where a button needs to be pressed, the systems decide whether a reference containing the color of the button

and the direction to the button will uniquely identifies the target button. That is, whether there are distracting buttons of the same color in the same direction. If there are none, both systems simply generate a RE like this one: “Press the green button behind you.” However, if this (first) RE would not uniquely identify the button, additional information must be provided. The additional information picks out the target button from the group of distractors. That is done either through a landmark, which I have placed in worlds so it is possible in most cases or it is done in a similar way the door specifying RE were created. That is sorting the group of distractors and the target button by their distance to the IF and using the target button’s index. Both systems are flexible enough to be able to provide unique identification to a substantial number of cases.

The system Beta will present that additional information as part of the first RE. System Alpha, on the other hand, separates this additional information to a new RE, which is not presented immediately. The Alpha system waits until the target button is visible and only then presents the second RE. Since the first RE contains direction information, it is presumed the IF will start moving towards the target button.

An example of instructions by Alpha system is in the following text. Spatial information is transcribed in parentheses for the sake of clearness.

(IF enters a new room.)
Alpha: Press the red button on your left.
(IF turns left and sees three red buttons.)
Alpha: The one next to the flower.
(IF goes and press the correct button)

An example of instructions by Beta system is in the following text.

(IF enters a new room.)
Beta: Press the red button on your left, the one next to the flower.
(IF turns left and goes press the red button.)

In cases where additional information is needed, both systems provide positive feedback when the IF is close and looking at the target button, measured by multiplying the distance between the IF and target button and the angle between them and comparing it to a threshold. Once a button was pressed, either positive or negative feedback is generated, depending whether the IF pressed the correct button.

That concludes the high-level overview of the Alpha and Beta systems. I’ll now present the results of the experiment.

5.3 Results

I’ve recruited X colleagues and friends, mostly university students. They all had at least basic knowledge of English language to understand the instructions. The participants together played Z valid games on 4 evaluation worlds.

Overall Y cases, where a target button has to be specified, were recorded. Of those A were handled by Alpha system and B were handled by Beta system. The means and standard deviations for the two systems can be found in table 5.1.

System	Average (s)	SD (s)
Alpha	$m1$	$std1$
Beta	$m2$	$std2$

Table 5.1: Results of the experiment.

The t-test statistics for independent two-tailed t-test is Z with p-value of p . Therefore there isn't sufficient evidence to refuse the null hypothesis that the mean times between a first reference and the button press are equal for Alpha and Beta systems.

This experiment was only testing very specific reference splitting. The first reference always contained the position and color of the button. In case of system Beta it also contained relative position in a group of distractors or to a landmark. For system Alpha that last specification reference was delayed until the target button was visible. Because of the written instruction, the splitting system Alpha had a small disadvantage, since the second reference in some cases had to wait for the first one to finish, which slowed the referencing down. It would be interesting to explore the splitting in other scenarios and possibly in spoken navigation. But that was beyond scope of this thesis.

Based on the result, I would argue that the strategy of splitting reference is a language device to reduce complexity of REs in complex environments, which however does not significantly affects the task performance. It is an interesting phenomenon which I believe would be worth exploring for REG research.

Conclusion

This thesis analysed a dataset of human to human interaction in a shared navigation task. Focusing on referring expressions, it revealed a new referencing strategy, which is not following the methodology of previous research in REG field.

After describing the shared task and the dataset, this thesis attempted to model the referencing strategy and other related NLG problems using machine learning techniques. Spatial features, such as information about the scene complexity, were successfully used in related research and were therefore extracted from the dataset. They, however, proved to have less predictive power in this task. This paper argues that the complexity of the shared task and dataset size had an influence on these results.

Finally, this thesis talks about a conducted experiment on the mentioned referencing strategy. The task proficiency of two NLG systems was compared in this experiment. One of the systems represented the newly discovered strategy, while the other one followed more standard methodology of REG. No significant difference between the systems' proficiency was found, suggesting the new referencing strategy neither improves nor harms system's efficiency and has more to do with human preferences and personal strategies.

It would be interesting to extract features which include more information about instruction givers and followers from the dataset and see if the attempted machine learning techniques will improve. Focusing on local and personal referencing strategies rather than global ones, is another way to continue in the work of this thesis. It might also be worth exploring a different dataset with similar complexity. For future studies I would suggest exploring splitting references into multiple REs as an interesting strategy to explore.

Bibliography

- Herbert H Clark and Meredyth A Krych. Speaking while monitoring addressees for understanding. *Journal of Memory and Language*, 50(1):62–81, 2004.
- Robert Dale and Ehud Reiter. Computational interpretations of the gricean maxims in the generation of referring expressions. *Cognitive Science*, 19(2):233–263, 1995.
- Carlos Gómez Gallo, T Florian Jaeger, James F Allen, and Mary D Swift. Production in a multimodal corpus: how speakers communicate complex actions. In *LREC*, 2008.
- Andrew Gargett, Konstantina Garoufi, Alexander Koller, and Kristina Striegnitz. The give-2 corpus of giving instructions in virtual environments. In *LREC*, 2010.
- David C Geary, Scott J Saults, Fan Liu, and Mary K Hoard. Sex differences in spatial cognition, computational fluency, and arithmetical reasoning. *Journal of Experimental Child Psychology*, 77(4):337–353, 2000.
- Eun Young Ha, Joseph F Grafsgaard, Christopher M Mitchell, Kristy Elizabeth Boyer, and James C Lester. Combining verbal and nonverbal features to overcome the ‘information gap’ in task-oriented dialogue. In *Proceedings of the 13th Annual Meeting of the Special Interest Group on Discourse and Dialogue*, pages 247–256. Association for Computational Linguistics, 2012.
- Mark Hall, Eibe Frank, Geoffrey Holmes, Bernhard Pfahringer, Peter Reutemann, and Ian H Witten. The weka data mining software: an update. *ACM SIGKDD Explorations Newsletter*, 11(1):10–18, 2009.
- Mary Hegarty, Anthony E Richardson, Daniel R Montello, Kristin Lovelace, and Ilavanil Subbiah. Development of a self-report measure of environmental spatial ability. *Intelligence*, 30(5):425–447, 2002.
- Rodger Kibble and Richard Power. An integrated framework for text planning and pronominalisation. In *Proceedings of the first international conference on Natural language generation-Volume 14*, pages 77–84. Association for Computational Linguistics, 2000.
- Alexander Koller, Kristina Striegnitz, Donna Byron, Justine Cassell, Robert Dale, Johanna Moore, and Jon Oberlander. The first challenge on generating instructions in virtual environments. In *Empirical Methods in Natural Language Generation*, pages 328–352. Springer, 2010a.

- Alexander Koller, Kristina Striegnitz, Andrew Gargett, Donna Byron, Justine Cassell, Robert Dale, Johanna Moore, and Jon Oberlander. Report on the second nlg challenge on generating instructions in virtual environments (give-2). In *Proceedings of the 6th international natural language generation conference*, pages 243–250. Association for Computational Linguistics, 2010b.
- Alexander Koller, Maria Staudte, Konstantina Garoufi, and Matthew Crocker. Enhancing referential success by tracking hearer gaze. In *Proceedings of the 13th Annual Meeting of the Special Interest Group on Discourse and Dialogue*, pages 30–39. Association for Computational Linguistics, 2012.
- Emiel Krahmer and Kees Van Deemter. Computational generation of referring expressions: A survey. *Computational Linguistics*, 38(1):173–218, 2012.
- Max M Louwerse, Nick Benesh, Mohammed E Hoque, Patrick Jeuniaux, Gwyneth Lewis, Jie Wu, and Megan Zirnsstein. Multimodal communication in face-to-face computer-mediated conversations. In *Proceedings of the 28th Annual Conference of the Cognitive Science Society*, pages 1235–1240, 2007.
- Scott D Moffat, Elizabeth Hampson, and Maria Hatzipantelis. Navigation in a "virtual" maze: Sex differences and correlation with psychometric measures of spatial ability in humans. *Evolution and Human Behavior*, 19(2):73–87, 1998.
- F. Pedregosa, G. Varoquaux, A. Gramfort, V. Michel, B. Thirion, O. Grisel, M. Blondel, P. Prettenhofer, R. Weiss, V. Dubourg, J. Vanderplas, A. Passos, D. Cournapeau, M. Brucher, M. Perrot, and E. Duchesnay. Scikit-learn: Machine learning in Python. *Journal of Machine Learning Research*, 12:2825–2830, 2011.
- Han Sloetjes and Peter Wittenburg. Annotation by category: Elan and iso dcr. In *LREC*, 2008.
- Laura Stoia, Donna K Byron, Darla Magdalene Shockley, and Eric Fosler-Lussier. Sentence planning for realtime navigational instructions. In *Proceedings of the Human Language Technology Conference of the NAACL, Companion Volume: Short Papers*, pages 157–160. Association for Computational Linguistics, 2006a.
- Laura Stoia, Darla Magdalene Shockley, Donna K Byron, and Eric Fosler-Lussier. Noun phrase generation for situated dialogs. In *Proceedings of the fourth international natural language generation conference*, pages 81–88. Association for Computational Linguistics, 2006b.
- Kristina Striegnitz, Alexandre Denis, Andrew Gargett, Konstantina Garoufi, Alexander Koller, and Mariët Theune. Report on the second second challenge on generating instructions in virtual environments (give-2.5). In *Proceedings of the 13th European Workshop on Natural Language Generation*, pages 270–279. Association for Computational Linguistics, 2011.
- Kristina Striegnitz, Hendrik Buschmeier, and Stefan Kopp. Referring in installments: a corpus study of spoken object references in an interactive virtual environment. In *Proceedings of the Seventh International Natural Language Generation Conference*, pages 12–16. Association for Computational Linguistics, 2012.

Jette Viethen, Robert Dale, and Markus Guhe. Generating subsequent reference in shared visual scenes: Computation vs. re-use. In *Proceedings of the Conference on Empirical Methods in Natural Language Processing*, pages 1158–1167. Association for Computational Linguistics, 2011a.

Jette Viethen, Robert Dale, and Markus Guhe. The impact of visual context on the content of referring expressions. In *Proceedings of the 13th European Workshop on Natural Language Generation*, pages 44–52. Association for Computational Linguistics, 2011b.

Appendix

.1 Histograms for first reference ML

This section contains attributes' histograms for timing of the first reference ML.




Figure 5: Histogram of attribute distance to the target

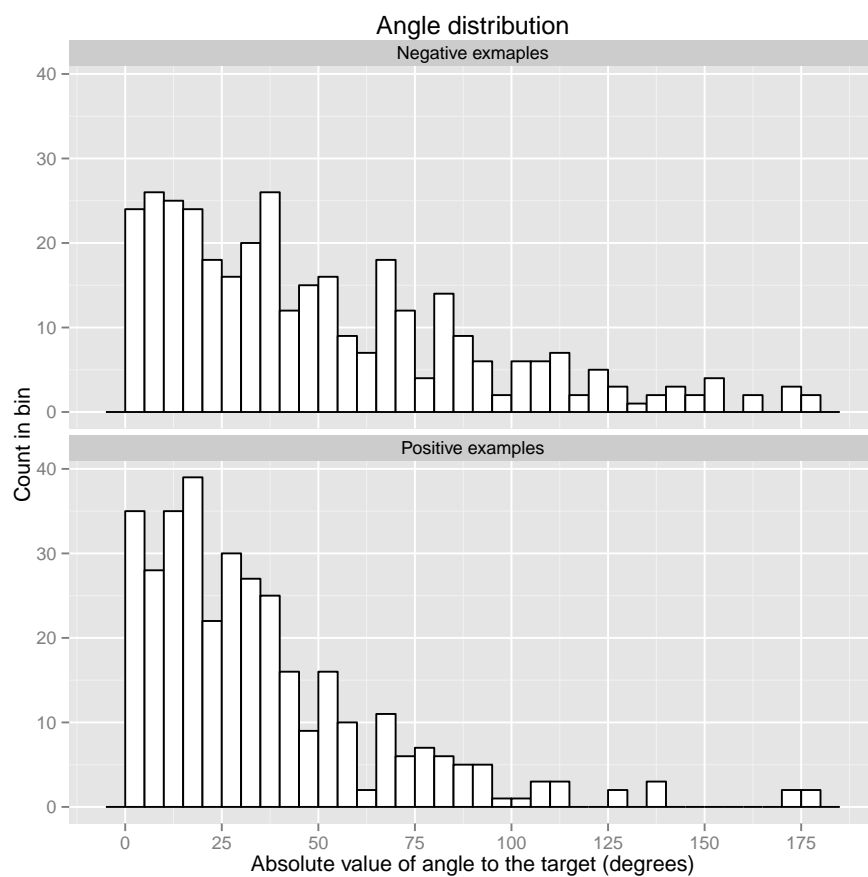


Figure 6: Histogram of attribute angle to the target




Figure 7: Histogram of attribute whether the target is visible




Figure 8: Histogram of attribute number of distractors




Figure 9: Histogram of attribute number of distracting buttons




Figure 10: Histogram of attribute number of distracting buttons with lesser angle

.2 Histograms for chains ML




Figure 11: Histogram of attribute time spent not moving




Figure 12: Histogram of attribute time spent rotating




Figure 13: Histogram of attribute objects in the room




Figure 14: Histogram of attribute buttons in the room




Figure 15: Histogram of attribute landmarks in the room




Figure 16: Histogram of attribute very close buttons to the target




Figure 17: Histogram of attribute buttons on the same wall as the target




Figure 18: Histogram of attribute close buttons to the target




Figure 19: Histogram of attribute far buttons from the target